



[◀ Return to "Deep Reinforcement Learning Nanodegree" in the classroom](#)

[DISCUSS ON STUDENT HUB](#)

Navigation

REVIEW

CODE REVIEW

HISTORY

Meets Specifications

Dear Udacian, the project is very well implemented and meets the specifications! Congratulations on successfully completing the project.

As next step, please go through the resource: [human-like robot hand trained to manipulate physical objects with unprecedented dexterity](#).

All the best for the next projects! 😊

Training Code

The repository (or zip file) includes functional, well-documented, and organized code for training the agent.

The repository contains jupyter notebook, code files, readme, and project report. The code is functional. The implemented Deep Q Network and Double Deep Q Network look good!

You should definitely check the following resources to progress with the learning further:

- [Speeding up DQN on PyTorch: how to solve Pong in 30 minutes](#)
- [Advanced DQNs: Playing Pac-man with Deep Reinforcement Learning by mapping pixel images to Q](#)

values

The code is written in PyTorch and Python 3.

The code is written in pytorch and python3.

- A good read: [PyTorch vs TensorFlow—spotting the difference](#)

The submission includes the saved model weights of the successful agent.

Thanks for including the saved model weights of the successful agent.

Check points are very important, a good article - [Checkpointing Tutorial for TensorFlow, Keras, and PyTorch](#).

README

The GitHub (or zip file) submission includes a `README.md` file in the root of the repository.

The README describes the the project environment details (i.e., the state and action spaces, and when the environment is considered solved).

Awesome work providing the project environment details in the README.

State space, action space, reward function and when the environment is considered solved is specified very informatively.

The README has instructions for installing dependencies or downloading needed files.

Proper instructions have been specified in the README to install the dependencies and download the necessary files.

The README describes how to run the code in the repository, to train the agent. For additional resources on creating READMEs or using Markdown, see [here](#) and [here](#).

Report

The submission includes a file in the root of the GitHub repository or zip file (one of `Report.md`, `Report.ipynb`, or `Report.pdf`) that provides a description of the implementation.

Report has been included in the root of the github repository.

The report clearly describes the learning algorithm, along with the chosen hyperparameters. It also describes the model architectures for any neural networks.

The report is rather informative providing an insight on every aspect of the project which includes implementation, model architectures, hyperparameters, rewards, future works.

- Good implementation of the agents for Deep Q Network and Double Deep Q Network.
- Correct decoupling of the parameters being updated from the ones that are using a target network to produce target values.
- Perfect implementation of The Epsilon-greedy action selection to encourage exploratory behavior in the agent.
- Good use of tau parameter to perform soft-update. It helps to prevent variance into the process due to individual batches.
- Good use of the replay memory to store and recall experience tuples.
- The implementation is easy to debug and easily extensible.

A plot of rewards per episode is included to illustrate that the agent is able to receive an average reward (over 100 episodes) of at least +13. The submission reports the number of episodes needed to solve the environment.

Awesome

- The agent seems to perform very well!
- The agent is able to achieve an average score of 13+ over last 100 episodes in just
 - 198 episodes for DQN!
 - 183 episodes for DDQN!
- The submission discusses the rewards plot obtained clearly.

The submission has concrete future ideas for improving the agent's performance.

Great job providing the ideas to experiment more in future with the project!

I would like to point you to the following resources:

- [Rainbow: Combining Improvements in Deep Reinforcement Learning](#)
- [Conquering OpenAI Retro Contest 2: Demystifying Rainbow Baseline](#)

You should try implementing Prioritized Experience Replay also. It helps to improve the performance and significantly reduces the training time. A fast implementation of Prioritized Experience Replay is possible using a special data structure called a Sum Tree. I found a [good implementation here](#).

And, you should definitely try applying these algorithms by taking raw screen pixels as input also. In case you get stuck, definitely check [this github repository](#).

 [DOWNLOAD PROJECT](#)

[RETURN TO PATH](#)